

Guess-and-Learn (G&L): Measuring the Cumulative Error Cost of Cold-Start Adaptation

Roland Arnold
rolandwarnold@gmail.com

Abstract

Evaluation of machine learning models typically emphasizes final accuracy, overlooking the cost of adaptation: the cumulative errors incurred while learning from scratch. Guess-and-Learn (G&L) v1.0 addresses this gap by measuring cold-start adaptability—the total mistakes a model makes while sequentially labeling an unlabeled dataset. At each step, the learner selects an instance, predicts its label, receives the ground truth, and updates parameters under either online (per-sample) or batch (delayed) mode. The resulting error trajectory exposes adaptation speed, selection quality, and bias—dynamics invisible to endpoint metrics.

G&L defines four tracks (Scratch/Pretrained \times Online/Batch) to disentangle the effects of initialization and update frequency. We formalize the protocol, relate it to classical mistake-bound theory, and estimate a heuristic “oracle reference band” for MNIST as a plausibility reference. Baseline experiments on MNIST and AG News, spanning classical methods (Perceptron, k -NN), convolutional architectures (CNN, ResNet-50), and pretrained transformers (ViT-B/16, BERT-base), reveal systematic differences in early-phase efficiency: smaller models can adapt with fewer initial errors, while pretraining benefits vary by domain. Across settings, current models remain well above the oracle band, highlighting an adaptability gap.

By quantifying the mistake cost of early learning, G&L complements conventional benchmarks and provides a reproducible framework for developing learners that are not only accurate in the limit but also reliable from the first examples.

1 Introduction

Modern machine learning systems achieve high accuracy, but this success often depends more on scale than adaptability. Most models require large amounts of labeled data, learn slowly from scratch, and adapt poorly to new domains where in-domain labels are scarce. In contrast, truly efficient learners would achieve competence with fewer examples and recover more quickly under distribution shift.

Evaluation practice, however, is dominated by two metrics: final accuracy on a held-out test set and label efficiency, measuring how many labeled samples are needed to reach a target performance. While useful, both overlook the cost of adaptation—the total mistakes a model makes while learning. Two models may reach similar accuracy but differ drastically in how many errors they incur along the way.

Guess-and-Learn (G&L) addresses this gap by treating cumulative error, not label count, as the primary metric. At each step, a learner selects an instance, predicts its label, receives the ground truth, and updates parameters. The benchmark tracks cumulative mistakes over the full labeling process, producing an error trajectory that directly reflects adaptation speed and efficiency. This perspective is especially relevant in domains where early mistakes carry real costs: degraded user experience in interactive systems [1], increased safety risk in autonomous systems, or wasted resources when labels are expensive to obtain.

Our central question is: *How many total errors does a model make when labeling every instance in a dataset, starting from zero in-domain knowledge?* Formally, at step t , the cumulative error is

$$E_t = \sum_{i=1}^t \mathbf{1}\{\hat{y}_i \neq y_i\}, \quad (1)$$

where \hat{y}_i is the model’s prediction, y_i the ground truth, and $\mathbf{1}\{\cdot\}$ the indicator function.

For humans—or an idealized clustering-and-reasoning oracle—the answer can be remarkably low. On MNIST, probabilistic analysis [25] suggests a plausible floor in the single digits (see Appendix C), and under realistic assumptions accounting for ambiguous instances [36], mastery is estimated to be possible after fewer than twenty errors. In practice, even state-of-the-art models make tens or hundreds of times more mistakes, underscoring a clear adaptability gap.

G&L is not a replacement for existing benchmarks, but a complementary diagnostic. By quantifying the mistake cost of early learning, it provides a concrete, reproducible metric for assessing model adaptability and for driving the development of learners that are not only accurate in the limit but also reliable from the very first examples.

Contributions. This paper’s contributions are:

- **A new evaluation protocol (G&L)** for measuring cold-start adaptation, defined by four tracks that isolate initialization and update frequency.
- **A heuristic oracle reference band** for MNIST to provide a plausible performance floor for early-phase errors.
- **A reproducible experimental framework**, including open-source code with a pinned commit and fixed seeds.
- **Benchmarks of classic and pretrained models** under five acquisition strategies, reporting their cumulative-error trajectories.
- **An analysis of training policy effects** (e.g., update cadence, weight resets) that highlights a capacity-agility trade-off in early adaptation.

2 Purpose & Scope

The goal of Guess-and-Learn (G&L) is not simply to produce another leaderboard, but to establish a benchmark that makes early-phase adaptability measurable. By quantifying the mistakes made during cold-start learning, G&L turns variation in efficiency into a concrete, comparable target for model improvement.

Cold-start setting. In G&L, the learner begins with:

- Zero in-domain labeled instances.
- Optional task-agnostic prior knowledge (e.g., pretrained weights).
- A requirement to predict a label before receiving the ground truth.
- No option to abstain from prediction.
- Freedom to select the next instance to label by any strategy.

These conditions hold throughout. Any form of manual seeding, few-shot prompting, or oracle hints constitutes a warm start and is outside the scope of G&L v1.0.

Protocol. The protocol involves a sequential, full-pool labeling process. At each step, the learner (i) selects an unlabeled example from the pool, (ii) predicts its class, (iii) receives the ground-truth label from an oracle, and (iv) updates its parameters based on a fixed schedule. Every prediction contributes to cumulative error. The primary performance measure is the final cumulative error, E_N , after labeling a pool of N instances.

Tracks. G&L v1.0 defines four tracks that isolate the influence of initialization and update frequency:

- **Scratch–Online (SO):** random initialization; update after every labeled example.
- **Scratch–Batch (SB):** random initialization; update after batches of size $K > 1$.
- **Pretrained–Online (PO):** initialized from pretrained weights; update after every labeled example.
- **Pretrained–Batch (PB):** initialized from pretrained weights; update after batches of size $K > 1$.

These tracks separate the influence of prior knowledge (scratch vs. pretrained) from the influence of update schedule (per-step vs. batch).

Oracle reference band (MNIST). To contextualize MNIST results, we include a heuristic oracle band: an estimate of the minimum plausible number of errors under idealized assumptions (e.g., perfect clustering with immediate generalization after one labeled example per class). Its width reflects uncertainty from factors such as class imbalance, intra-class variation, and label noise [27]. This heuristic estimate is not a formal bound but an illustrative reference, included to highlight the large gap between current systems and efficient learning under ideal conditions (Appendix C).

3 Relation to Existing Paradigms

Guess-and-Learn (G&L) borrows familiar settings from existing paradigms but alters the objective to measure the cumulative error cost of cold-start adaptation.

Active learning. Pool-based active learning seeks to minimize label queries needed to reach a target accuracy [8, 32]. G&L uses the same interactive pool-selection setting but inverts the objective: it measures the total mistakes made while labeling the entire pool. Two reference points are useful: (i) a random baseline: for C balanced classes, uniform guessing incurs $\mathbb{E}[E_N] = N(1 - 1/C)$; and (ii) an existence floor: at least $C - 1$ mistakes are unavoidable in the worst case, since the learner must encounter at least one example of each class before knowing it exists [2].

Online learning. The G&L Scratch–Online (SO) track is a direct empirical analogue of the mistake-bound model [22]. For linearly separable data with margin γ and diameter R , the Perceptron has the classic bound $E_N \leq R^2/\gamma^2$ [28, 22]. G&L-SO instantiates this setting on modern, high-dimensional data, with the added twist that the learner actively selects the sequence of examples.

KWIK learning (abstention). “Knows What It Knows” frameworks allow abstention until confidence is high, thereby avoiding mistakes [21]. G&L forbids abstention; every instance

must be predicted. This models mandatory-decision scenarios (e.g., medical triage, autonomous navigation) and ensures the full cost of uncertainty is counted.

Curriculum & self-paced learning. Acquisition strategies impose a curriculum [3, 18]. Uncertainty sampling tends toward “hard-first”, while confidence sampling is “easy-first”. Heuristically, if a strategy maintained per-step margin $\gamma_t \geq \gamma_{\min}$, the Perceptron bound tightens to

$$E_N \leq \frac{R^2}{\gamma_{\min}^2}$$

This motivates margin-aware sampling, though it is not a guarantee for the heuristics we evaluate.

Prequential evaluation. G&L’s predict-then-update loop mirrors prequential evaluation in data streams [10, 16], with three differences: (i) a fixed pool in cold-start; (ii) cumulative error is the primary metric; and (iii) the learner controls instance order.

4 Methodology

4.1 G&L Protocol

At each step the learner:

- (i) selects one unlabeled instance using a chosen acquisition strategy;
- (ii) predicts its label;
- (iii) receives the ground-truth label from an oracle;
- (iv) updates its parameters according to a predefined schedule.

Predictions are mandatory (no abstention), and errors are accumulated over time. The primary metric is final cumulative error E_N after all N instances are labeled.

4.2 Benchmark Tracks

We cross two factors—initialization and update schedule—to form four tracks (Table 1). Online tracks update after each instance; batch tracks update after groups of size $K > 1$ (default $K = 50$ unless otherwise specified).

Table 1: G&L tracks formed by the cross of initialization and update schedule.

Track	Initialization	Update Schedule	Batch Size (K)
G&L-SO	Scratch	Online	1
G&L-SB	Scratch	Batch	> 1 (default $K=50$)
G&L-PO	Pretrained	Online	1
G&L-PB	Pretrained	Batch	> 1 (default $K=50$)

4.3 Acquisition Strategies

We evaluate five established strategies:

- **Random:** uniform sampling without replacement.

- **Confidence (easy-first):** select x maximizing $\max_y P_\theta(y | x)$ [3].
- **Least-confidence:** select x that minimizes $\max_y P_\theta(y | x)$ [20].
- **Margin:** select x with smallest gap $p_1 - p_2$ between the top two class probabilities [31, 34].
- **Entropy:** select x with largest $-\sum_y P_\theta(y | x) \log P_\theta(y | x)$ [32].

These are interpretable baselines to reveal how acquisition logic shapes the error trajectory; we do not claim optimality for G&L.

4.4 Implementation Details

All models produce class-probability vectors $P_\theta(y | x)$ with prediction $\hat{y} = \arg \max_y P_\theta(y | x)$. Models producing logits use a softmax [6]; k -NN probabilities derive from distance-weighted neighbor votes (normalized) [9]. For k -NN, 'learning' consists of adding each newly labeled instance to the pool of reference examples used for subsequent distance-weighted predictions. Ties (in acquisition or prediction) are broken uniformly at random. The pool is processed without replacement until exhaustion. Unless noted, results are averaged over fixed seeds.

4.5 Experimental Setup

Datasets. MNIST (image classification) [19] and AG News (text classification) [37]. For G&L, the unlabeled pool is drawn from each dataset's official test set and processed without replacement, solely to standardize the unlabeled pool for measuring adaptation cost; we do not report conventional test accuracy here, so this does not create leakage for our target metric.

Models. *Vision (MNIST):* k -NN ($k=7$) [9, 14], Perceptron [29], a small 3-layer CNN [19], ResNet-50 [17], ViT-B/16 [12, 35]. *Text (AG News):* k -NN, linear Perceptron, BERT-base [11, 35].

Pretrained tracks. In PO (Pretrained–Online), pretrained backbones are used as frozen feature extractors with a newly initialized linear head trained online. In PB (Pretrained–Batch), all parameters are fine-tuned end-to-end on the batch schedule.

Acquisition. All five strategies from 4.3 are used.

Updates. Online tracks update every step; batch tracks use $K > 1$ (default $K=50$). The ablation varies $K \in \{10, 50, 200\}$.

Reporting. Unless noted, results are averaged over seeds and reported for $n=300$ early-phase subsets; full-pool curves are included for MNIST to contextualize the oracle band.

5 Results and Analysis

We organize findings into four themes: model adaptability, strategy effects, training policies, and cost–performance trade-offs.

5.1 The Adaptability Gap

Most models are far from optimal in early adaptation. On MNIST, the best-performing ViT-B/16 track—pretrained-batch (PB)—accumulated 136.0 ± 4.75 errors at $n = 300$ under Random acquisition (mean \pm standard error across seeds), a finding visualized in Figure 1. In contrast, a conceptual clustering oracle—requiring only one representative example per class—suggests a 7–12 mistake plausibility band under idealized assumptions (see Appendix C). For AG News, whose class boundaries are more ambiguous (e.g., *World* vs. *Business*), and where prior work reports nontrivial disagreement and label noise [27], we do not provide an oracle estimate.




Figure 1: MNIST cumulative error trajectories with the heuristic oracle band (7–12). All models remain well above the band during early adaptation.

5.2 Capacity vs. Agility

Adaptability is not a simple function of size. In cold-start, smaller online learners can outperform larger models early on (Figure 2). On the first 300 MNIST samples, the Perceptron incurs fewer cumulative errors than both a CNN and k -NN. High capacity (ResNet-50, ViT-B/16) does not immediately translate into efficient boundary formation when labeled data are scarce.

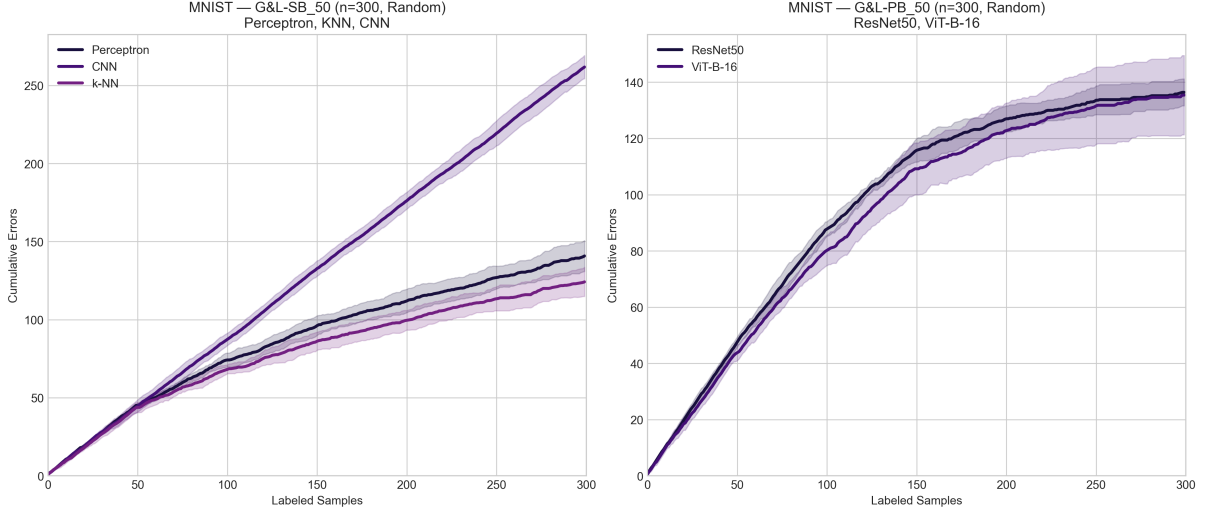


Figure 2: Capacity vs. agility on MNIST ($n=300$, Random acquisition). Left: small SO models (Perceptron, k -NN, CNN). Right: pretrained PB₅₀ models (ResNet-50, ViT-B/16). Smaller online learners adapt faster.

5.3 Impact of Acquisition Strategies

Effects are task- and model-dependent, as shown in Figure 3. On MNIST with a CNN (SO), *Confidence* (easy-first) yields the lowest cumulative error over the first 300 samples. Uncertainty-based strategies (Entropy, Margin, Least-Confidence) underperform even Random sampling, likely because they front-load ambiguous cases that inflate early mistakes. On AG News with BERT-base (PO), strategy differences are small and often within variance bands in the first 300 samples; strong priors dominate.

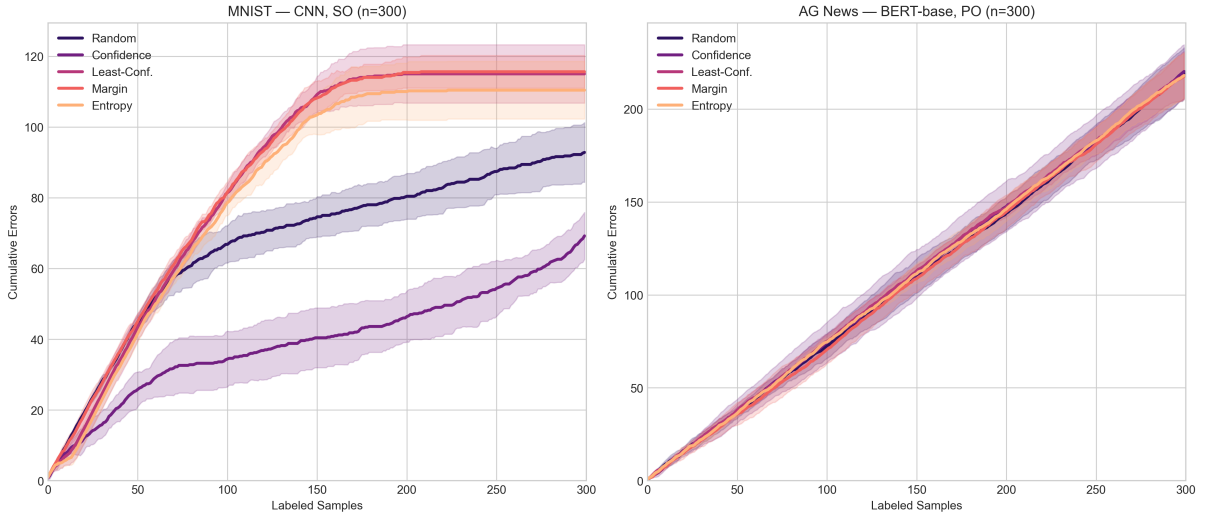


Figure 3: Strategy effects. Left: MNIST—CNN, SO ($n=300$): Confidence (easy-first) yields the lowest cumulative error; entropy/margin/least-confidence are worse than Random. Right: AG News—BERT-base, PO ($n=300$): curves are tightly clustered; strategy choice has minimal impact over the first 300 samples.

5.4 The Role of Training Policies

Two policy choices directly shape adaptability.

Update cadence (K). On MNIST (CNN, SB), more frequent updates ($K=10$) generally reduce errors, but the relationship is not monotonic: $K=200$ unexpectedly outperforms $K=50$, suggesting optimizer/stability effects.

Weight reset. On AG News (BERT, PB), resetting to the original pretrained weights before each batch increased errors; carrying forward the fine-tuned weights reduced cumulative error.




Figure 4: Ablations. Left: An ablation on the batch size K for a CNN on MNIST (SB track, Entropy acquisition). The impact of update frequency is non-monotonic; while more frequent updates ($K = 10$) are most effective, a very large batch size ($K = 200$) unexpectedly outperforms an intermediate one ($K = 50$). Right: AG News BERT-base, PB_{50} Entropy (resetting between batches increases errors).

5.5 Cost–Performance Trade-offs

Plotting final error against mean wall-clock time ($n=300$) reveals distinct cost–performance frontiers by dataset and track. Each point plots the mean performance of a specific model–and–strategy combination; plots are faceted to avoid mixing regimes.

On MNIST, SO models define the most efficient frontier. k -NN achieves the lowest error (~ 80 – 100) at moderate cost, while the Perceptron is faster but higher-error; the CNN performs poorly in SO. Deferring updates in SB degrades performance. Pretrained vision models (ResNet-50, ViT-B/16) in PO/PB are costly and not competitive on this dataset.

On AG News, pretrained models are essential. In scratch tracks, Text k -NN is a reasonable baseline and outperforms the fast but error-prone Text Perceptron. BERT-base in PB (fine-tuning) attains the lowest errors (~ 80 – 130) at the highest cost; PO (frozen features) is weaker. Aggregating across strategies shifts frontiers (e.g., SO on MNIST toward k -NN).




Figure 5: Cost–performance ($n=300$): mean wall-clock time (log-scale) vs. mean final error, faceted by dataset (MNIST, AG News) and track (SO, SB, PO, PB). Each point represents the mean performance of a specific model–and–strategy pairing.

6 Discussion

6.1 Key Findings and Interpretations

Adaptability remains challenging. Even powerful pretrained models accumulate hundreds of early mistakes on simple tasks and remain far above the plausible 7–12 error band on MNIST.

Agility can outweigh capacity. On simpler tasks and in early stages, smaller online models (e.g., Perceptron) often incur fewer errors than high-capacity architectures, underscoring a capacity–agility tension.

Strategy is context-dependent. Uncertainty-based methods can hurt when they front-load ambiguous cases (e.g., MNIST with CNN), yet matter little when strong priors dominate (e.g., BERT on AG News).

Training policies are critical. Update frequency and weight retention substantially affect cumulative error; retaining learned information between batches is especially important.

Cost–performance is regime-specific. No single model excels across all tracks. The value of pretraining and sensitivity to batch updates are task dependent, arguing for regime-aware evaluation.

6.2 From Diagnosis to Design

The diagnostics point to concrete opportunities:

- **Architectures & optimizers** designed for rapid, stable per-sample or small-batch updates.
- **Transfer & meta-learning** aimed at improving cold-start error efficiency, not only final accuracy.
- **Hybrid update schemes** that leverage large-model capacity without sacrificing online agility.
- **Acquisition–update co-design** to optimize instance selection jointly with update dynamics.

6.3 Limitations and Future Work

G&L v1.0 prioritizes a clear, reproducible protocol, which entails scope limits and natural directions:

- **Dataset scope:** two well-curated benchmarks. Larger/noisier data may require streaming or subsampling variants.
- **Generalization:** the protocol measures adaptation on a fixed pool; it does not assess post-adaptation performance in new domains.
- **Excluded scenarios:** class imbalance, open-set recognition, abstention, and instance-dependent error costs are out of scope but important for future versions (e.g., reject option, ambiguity-aware acquisition, cost-weighted errors).

7 Conclusion

Guess-and-Learn addresses a clear gap by measuring the cumulative error cost of learning from a cold-start. By tracking mistakes step-by-step, it captures how quickly a model becomes useful—information that final accuracy and label-efficiency do not provide.

The four-track design (scratch vs. pretrained, online vs. batch) and full error trajectories expose actionable differences in adaptation: sensitivity to initialization and update policy, the variable impact of acquisition strategies, and regime-specific cost–performance trade-offs. These distinctions matter in practice: reducing early errors lowers annotation cost, mitigates risk in safety-critical settings, and improves user experience.

At the same time, results show a persistent adaptability gap. On MNIST, modern models make many more early mistakes than a plausible oracle. Closing this gap requires optimizing for agility as well as for final accuracy. We offer G&L not only as a benchmark but as a design driver: a concrete target and a reproducible way to assess progress toward learners that adapt faster, make fewer mistakes, and become reliable sooner.

References

- [1] S. Amershi, M. Cakmak, W. B. Knox, and T. Kulesza (2014). Power to the People: The Role of Humans in Interactive Machine Learning. *AI Magazine* 35(4):105-120.
- [2] D. Angluin (1988). Queries and Concept Learning. *Machine Learning* 2(4):319-342.
- [3] Y. Bengio, J. Louradour, R. Collobert, and J. Weston (2009). Curriculum learning. In *Proceedings of the 26th Annual International Conference on Machine Learning* (pp. 41-48).
- [4] H. D. Block (1962). The perceptron: a model for brain functioning. I. *Reviews of Modern Physics*, 34(1), 123-135.
- [5] L. Bottou (2010). Large-Scale Machine Learning with Stochastic Gradient Descent. In *Proceedings of COMPSTAT 2010*, pp. 177-186.
- [6] J. S. Bridle (1990). Probabilistic interpretation of feed-forward classification network outputs, with relationships to statistical pattern recognition. In F. Fogelman Soulie & J. Hérault (Eds.), *Neurocomputing: Algorithms, Architectures and Applications* (pp. 227-236). Springer.
- [7] O. Chapelle, B. Schölkopf, and A. Zien (2006). *Semi-Supervised Learning*. MIT Press.

- [8] D. Cohn, L. Atlas, and R. Ladner (1994). Improving Generalization with Active Learning. *Machine Learning* 15(2):201-221.
- [9] T. M. Cover, & P. E. Hart (1967). Nearest neighbor pattern classification. *IEEE Transactions on Information Theory*, 13(1), 21-27.
- [10] A. P. Dawid (1984). Present position and potential developments: Some personal views: Statistical theory: The prequential approach. *Journal of the Royal Statistical Society: Series A (General)*, 147(2), 278-292.
- [11] J. Devlin, M. W. Chang, K. Lee, & K. Toutanova (2019). BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)* (pp. 4171-4186).
- [12] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, & N. Houlsby (2021). An image is worth 16x16 words: Transformers for image recognition at scale. In *International Conference on Learning Representations*.
- [13] J. A. Fails and D. R. Olsen Jr. (2003). Interactive Machine Learning. In *Proceedings of the 8th International Conference on Intelligent User Interfaces (IUI)*, pp. 39-45.
- [14] E. Fix, & J. L. Hodges Jr (1951). *Discriminatory analysis. Nonparametric discrimination: consistency properties* (Project Number 21-49-004, Report Number 4). USAF School of Aviation Medicine, Randolph Field, Tex.
- [15] E. W. Forgy (1965). Cluster analysis of multivariate data: efficiency versus interpretability of classifications. *Biometrics*, 21(3), 768-780.
- [16] J. Gama, R. Sebastião, and P. P. Rodrigues (2013). On Evaluating Stream Learning Algorithms. *Machine Learning* 90(3):317-346.
- [17] K. He, X. Zhang, S. Ren, & J. Sun (2016). Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition* (pp. 770-778).
- [18] M. P. Kumar, B. Packer, & D. Koller (2010). Self-paced learning for latent variable models. *Advances in Neural Information Processing Systems*, 23, 1189-1197.
- [19] Y. LeCun, L. Bottou, Y. Bengio, & P. Haffner (1998). Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11), 2278-2324.
- [20] D. D. Lewis, & W. A. Gale (1994). A sequential algorithm for training text classifiers. In *SIGIR'94: Proceedings of the 17th Annual International ACM SIGIR Conference on Research and Development in Information Retrieval* (pp. 3-12).
- [21] L. Li, M. L. Littman, & T. J. Walsh (2008). Knows what it knows: a framework for self-aware learning. In *Proceedings of the 25th International Conference on Machine Learning* (pp. 568-575).
- [22] N. Littlestone (1988). Learning Quickly When Irrelevant Attributes Abound: A New Linear-threshold Algorithm. *Machine Learning* 2(4):285-318.
- [23] S. P. Lloyd (1982). Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2), 129-137.

- [24] J. MacQueen (1967). Some methods for classification and analysis of multivariate observations. In *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability*, 1(14), 281-297.
- [25] M. Mitzenmacher, & E. Upfal (2017). *Probability and Computing: Randomization and Probabilistic Techniques in Algorithms and Data Analysis* (2nd ed.). Cambridge University Press.
- [26] R. Motwani, & P. Raghavan (1995). *Randomized Algorithms*. Cambridge University Press.
- [27] C. G. Northcutt, L. Jiang, & I. L. Chuang (2021). Confident learning: Estimating uncertainty in dataset labels. *Journal of Artificial Intelligence Research*, 70, 1373-1411.
- [28] A. B. Novikoff (1962). On Convergence Proofs for Perceptrons. In *Proceedings of the Symposium on the Mathematical Theory of Automata*, Vol. 12, pp. 615-622.
- [29] F. Rosenblatt (1958). The perceptron: a probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6), 386-408.
- [30] S. Ross, G. Gordon, and J. A. Bagnell (2011). A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning. In *Proceedings of the 14th International Conference on Artificial Intelligence and Statistics (AISTATS)*, PMLR 15:627-635.
- [31] A. I. Schein, & L. H. Ungar (2007). Active learning for logistic regression: an evaluation. *Machine Learning*, 68(3), 235-265.
- [32] B. Settles (2009). *Active learning literature survey* (Computer Sciences Technical Report 1648). University of Wisconsin-Madison.
- [33] B. Settles and M. Craven (2008). An Analysis of Active Learning Strategies for Sequence Labeling Tasks. In *Proceedings of EMNLP 2008*, pp. 1070-1079.
- [34] S. Tong and D. K. Koller (2001). Support Vector Machine Active Learning with Applications to Text Classification. *Journal of Machine Learning Research* 2:45-66.
- [35] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, & I. Polosukhin (2017). Attention is all you need. *Advances in Neural Information Processing Systems*, 30.
- [36] P. Yadav and L. Bottou (2019). Cold Case: The Lost MNIST Digits, arXiv:1905.00047.
- [37] X. Zhang, J. Zhao, & Y. LeCun (2015). Character-level convolutional networks for text classification. *Advances in Neural Information Processing Systems*, 28.

A Experimental Reproducibility

Environment. All experiments were run on a single NVIDIA A100 (40 GB) GPU, using PyTorch 2.0.1, Hugging Face Transformers 4.54.0, and Python 3.11.

Scale. The analysis is based on 493 distinct experiments, which produced over 497,300 data records.

Seeds. Random seeds $[0, 1, 2, 3, 4]$ were fixed across runs.

Code and data. All scripts, configuration files, and logged results are available at: <https://github.com/RolandWArnold/guess-and-learn-benchmark>. The exact artifacts correspond to git commit 975886d4086f6844fc0458774f3de4989089e7c1. The repository README.md provides step-by-step reproduction instructions.

B Supplemental Visualizations

Figure 6 Early-stage adaptability on AG News ($n=300$). Pretrained BERT shows a clear advantage over scratch models, reflecting the benefit of task-aligned priors.




Figure 6: AG News early-stage adaptability ($n=300$).

Figure 7 Early-stage adaptability on MNIST ($n=300$). The Perceptron adapts quickly and outperforms deeper pretrained models early on, illustrating the capacity–agility trade-off.




Figure 7: MNIST early-stage adaptability ($n=300$).

C Heuristic Oracle Floor for MNIST

We provide a practical, heuristic estimate of a plausible floor on errors for MNIST—the “oracle floor”—based on a simple *cluster-first* oracle and an audit of residual ambiguity. This is an illustrative reference, not a formal bound.

Cluster-First Oracle

Setup. (i) Embed the unlabeled pool; (ii) partition into $C=10$ clusters (e.g., k -means [24, 23, 15]); (iii) query one medoid per cluster and assign its label to the cluster.

Mapping cost. Associating C abstract clusters to C class names is a coupon-collector variant with feedback [25, 26]. The expected number of mapping errors is

$$\mathbb{E}[E_{\text{map}}] = C - H_C, \quad H_C = \sum_{j=1}^C \frac{1}{j}. \quad (2)$$

For MNIST ($C=10$), this gives $\mathbb{E}[E_{\text{map}}] \approx 10 - H_{10} \approx 7.07$.

Residual ambiguity. Even with near-perfect clustering, two sources remain:

- *Boundary cost* (E_{bound}): cluster-purity audits suggest ~ 2 – 4 borderline cases (e.g., 1 vs. 7, 0 vs. 6).
- *Noise/ambiguity cost* (E_{noise}): label audits identify ~ 4 – 6 ambiguous or mislabeled instances (e.g., via Cleanlab; see [27, 36]).

These overlap with one another and with mapping steps; they are not simply additive.

Oracle band. Combining the mapping expectation (~ 7 errors) with a small overlapping residual yields a MNIST oracle band of roughly 7–12 errors. This represents a minimum plausible range

under idealized conditions, with remaining mistakes attributable to intrinsic dataset ambiguity rather than learner deficiencies.